

Pediment Bank · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Pediment Bank dbt-wizard demo.

3-5 days

MANUAL MODEL BUILD
TODAY

90 sec

DBT-WIZARD BUILD

\$2.4M

CNP FRAUD BLOCKED / WK

8 hr

EARLIER THAN RISK +
FRAUD DAILY STANDUP

Frame the demo on business outcomes — never functionality

PAIN The fraud team can't size new spikes before they bleed	REASON Every new question = 3-5 day DE backlog	IMPACT Realized fraud losses, OCC findings, blocked-customer false positives	CAPABILITY Build-time AI on the governed gold layer	PROOF Live model build in 90 seconds — on screen
---	--	--	---	--

The mantra for this video: the win is not that dbt-wizard *can* author a model. The win is that the Fraud Operations Director walks into the next Risk + Fraud daily standup with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties back to that one sentence.

The 3-minute script · five beats

00:00 – 00:25 Beat 1 · The stakeholder hook

WHAT YOU SAY (VERBATIM, 25S)

"It is Friday morning at Pediment Bank. The Fraud Operations Director needs to know: Why did CNP fraud spike for tier-2 merchants in Texas last Tuesday? The Risk + Fraud daily standup happens in 8 hours. Today, an answer like this costs three-to-five days of data-engineering work and a missed meeting. We are about to change that."

WHAT YOU SHOW

Land on the **Scenario page** (</scenario>). The Fraud Operations Director's question banner is the first thing in frame. Let the T-minus counter tick.

OUTCOME → The audience anchors on a real business question, not a data feature.

00:25 – 00:50 Beat 2 · The cost of the gap

WHAT YOU SAY (25S)

"There is no gold-layer mart that answers this. Three days from now, after a request, a backlog, a scope, a build, and a peer review, the engineering team might land it. The Risk + Fraud daily standup will have already happened. A 2.4x tier-2 spike in Texas — every quiet day means about \$2.4 million in card-not-present exposure that nobody is acting on — and that is just the dollars we can see."

WHAT YOU SHOW

Stay on Scenario. Highlight the four KPI tiles: Manual ETA · 3-5 days, dbt-wizard ETA · 90 sec, \$2.4M at risk, T-minus to Risk + Fraud daily standup. Slowly hover the "Manual ETA" tile.

OUTCOME → Quantifies the business cost of the gap. Sets up the 90-second proof point.

00:50 – 02:20 Beat 3 · The live build (the proof moment)

WHAT YOU SAY (90S)

"We click *Run the Build*. Four sub-agents — *Explorer*, *Summary*, *Worker*, *Verification* — open the dbt project, find the right grain, pull a slice of the underlying gold layer, write the new model, write its tests, and materialize it. The viewer sees the SQL appear character by character. By the time we are done talking, the model exists, the tests have passed, and the lineage graph has updated."

While narrating, pause once at: "Worker is anchoring the CTE to the grain the Fraud Operations Director asked for — not a notebook-grade approximation." Then pause again at: "Verification is writing a uniqueness test on the natural key. This is a governed gold asset, not a SQL snippet pasted into Slack."

WHAT YOU SHOW

Click into [/wizard-live](#). Let the playback run at 1x — do not bump speed. The agents type the SQL and YAML in real time. Step rail moves through Discovery → Schema → Inspection → Model Creation → Test Authoring → Materialization. The terminal-window aesthetic looks like Claude Code.

OUTCOME → Audience sees the build-time AI work end-to-end. The "wow" is not the typing animation — it is the realization that the output is a tested, governed asset.

02:20 – 02:45 Beat 4 · The outcome (tie to the P&L)

WHAT YOU SAY (25S)

"The new mart is materialized. The Risk + Fraud daily standup convenes in 8 hours, not three days. The Fraud Operations Director walks in with the answer: three Houston-area merchant terminals chained to a single device-fingerprint cluster originated 67% of the spike — block at the terminal, not the cardholder. That is a merchant-side block, not a customer-friction change. And the next time anyone asks this question — card-present, ACH, wire, anywhere — the model is already there."

WHAT YOU SHOW

Cut to [/outcome](#). Show: materialized model card with row counts, tests passed, lineage updated, the root-cause insight panel, and the impact panel.

OUTCOME → Closes the loop from question to action. The model is reusable — first build pays for the second, third, tenth.

02:45 – 03:00 Beat 5 · The architectural ask

WHAT YOU SAY (15S)

"Fivetran lands every Pediment Bank source. dbt governs the layer. dbt-wizard authors the gaps. Snowflake reads it. That loop closes for fraud rings, AML alert triage, deposit-beta sensitivity, and OCC audit prep — all on the same governed lake. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."

WHAT YOU SHOW

Final card with the four-pillar loop, QR code to [fivetran.com/contact](#) or the booth landing page. Tag dbt Labs co-brand.

OUTCOME → Calls out the architectural pattern (not a feature). Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT IN THE DEMO
Risk standup sees the answer same-day	3-5 day DE backlog	"Acted on by tomorrow, not next week"
\$2.4M weekly CNP exposure sized in 90 seconds	Open exposure window	"The bleed stops on Tuesday, not Friday"
Reusable governed mart for every new spike	One-off Splunk-style scripts	"Houston this week, anywhere next week — same model"
Fraud team scales without DE headcount	Hiring DE backfill	"Your CDO stops being the bottleneck"
Examiner sees lineage + tests, not a query history	Manual model-risk evidence	"MRM-grade by default"

Anti-patterns · what NOT to say

- **Don't say:** "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety seconds from question to materialized answer."
- **Don't say:** "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake.
- **Don't say:** "dbt-wizard generates SQL." Say: "dbt-wizard authors a governed asset — schema contract, tests, lineage, ownership."
- **Don't say:** "Look at the typing animation." The animation is the texture, not the value.
- **Don't open with the architecture diagram.** Open with the Fraud Operations Director's question. Architecture is the closer.

Discovery prompts to anchor before the demo (if presenting live)

BEFORE YOU PRESS PLAY

- "When fraud spikes in a region you have not modeled before, how long until a new cut of data lands in the warehouse?"
- "How many of your fraud-team analysts can author governed SQL without a data-engineering ticket?"
- "What is the OCC's stance on shadow SQL queries in your fraud workflows today?"
- "Who is the Economic Buyer for new fraud-detection investment — CRO, CISO, or the CDO?"

AFTER THE DEMO LANDS

- "Which of those agent steps — discovery, schema, data inspection, model authoring, test authoring, materialization — is the one your team spends the most time on today?"
- "If we sized a pilot on three of your highest-value recurring stakeholder questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Likely audience: CRO, Head of Fraud, CISO, CDO."

Required follow-up materials

- One-pager · Pediment Bank · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/banking/
- Live demo URL · fivetran-jasonchletsos.github.io/Banking-ODI-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the Architecture tab of the demo
- Cross-industry credibility · Clarity Health (Healthcare flavor of the same loop)

Booth-floor reminder: the demo is built to run unattended for booth visitors. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.